Three rooflines: measured CPUs and theoretical H200 SXM

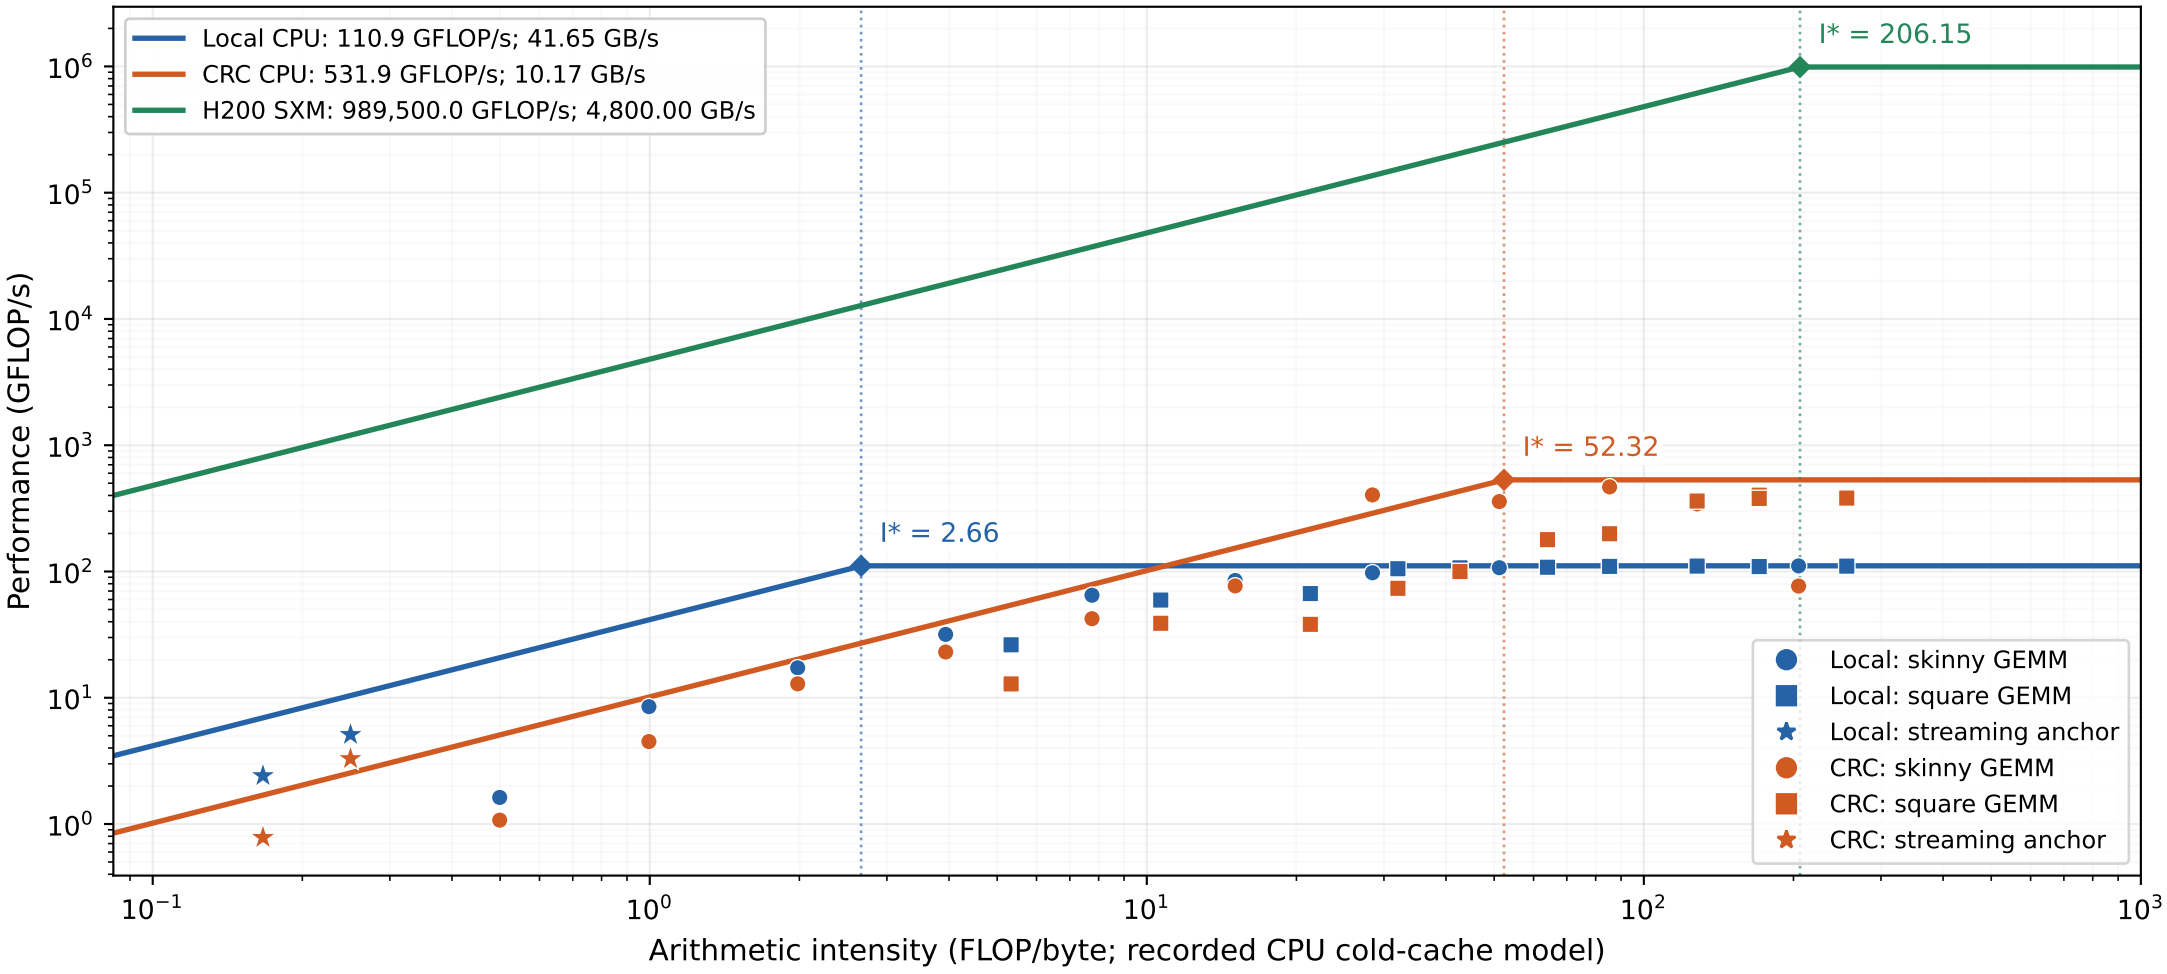

CPU points are measured FP32 and retain their original coordinates. CPU roofs use selected sustained calibrations.
H200 roof: one GPU, theoretical dense BF16/FP16 Tensor Core peak; no GPU performance was measured.
Precision changes arithmetic intensity: these CPU points are not BF16 performance predictions.

H200 source: https://www.nvidia.com/en-us/data-center/h200/
1,979 sparse TFLOP/s / 2 = 989.5 dense TFLOP/s; HBM bandwidth = 4.8 TB/s.